



Publish latency vs. batch fraction of capacity (median, shaded = 10th-90th percentile; dashed = extrapolation)

small (1 shard, log_cap=22)



medium (2 shards, log_cap=27)



large (128 shards, log_cap=27)

